

```

## Gitam Phase 1 Hackathon
## K Pranay 3061, A Sai Nikhil 2245, S Satish 4699
## MSC Data Science 'A'
## Model to Combat Cyberbullying.

## Importing required Python libraries and their functions.

import re
import os

import pandas as pd
import numpy as np
import string
from collections import Counter
import sklearn
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report

import tensorflow as tf
from tensorflow.keras import layers
from tensorflow.keras import losses
from tensorflow.keras import regularizers
from tensorflow.keras import preprocessing
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences

import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')

```

```

# Reading the raw data set
raw_data = pd.read_csv('/content/train_E6oV3lV.csv')
data = raw_data.copy()
data.drop(columns=['id'], axis=1, inplace=True)
data.head()

```

	label	tweet	
0	0	@user when a father is dysfunctional and is s...	
1	0	@user @user thanks for #lyft credit i can't us...	
2	0	bihday your majesty	
3	0	#model i love u take with u all the time in ...	
4	0	factsguide: society now #motivation	

Next steps:

[Generate code with data](#)[View recommended plots](#)

```

# Through this we found that how much is normal tweets and how much are of not relevant tweets

```

```

print(np.round(data['label'].value_counts()[0]/len(data) * 100, 2), "% are Normal Speech")
print(np.round(data['label'].value_counts()[1]/len(data) * 100, 2), "% are Hate Speech")

```

```

92.99 % are Normal Speech
7.01 % are Hate Speech

```

```

# Plotting it into the graphs

```

```

colors = ["#0101DF", "#DF0101"]

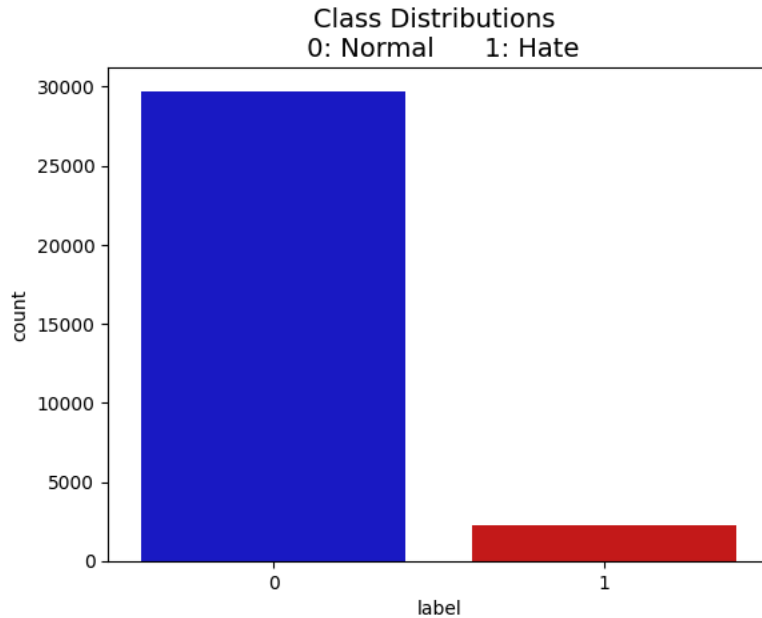
```

```

sns.countplot(x='label', data=data, palette=colors)
plt.title('Class Distributions \n 0: Normal      1: Hate', fontsize=14)

```

```
Text(0.5, 1.0, 'Class Distributions \n 0: Normal      1: Hate')
```



# Till now we have seen only for texts but emojis are also coming into bullying  
 # So for that we are now filtering the emojis

```
def remove_emoji(text):
    emoji_pattern = re.compile("["
        u"\U0001F600-\U0001F64F" #emoticons
        u"\U0001F300-\U0001F5FF" #symbols & pictograms
        u"\U0001F680-\U0001F6FF" #transport & map symbols
        u"\U0001F1E0-\U0001F1FF" #flags(ios)
        u"\U00002702-\U000027B0"
        u"\U000024C2-\U0001F251"
        "]" +, flags=re.UNICODE)
    return emoji_pattern.sub(r'', text)

def clean_text(text):
    delete_dict = {sp_character: '' for sp_character in string.punctuation}
    delete_dict[' '] = ' '
    table = str.maketrans(delete_dict)
    text1 = text.translate(table)
    textArr = text1.split()
    text2 = ' '.join([w for w in textArr if(not w.isdigit() and (not w.isdigit() and len(w) > 3))])

    return text2.lower()

# now we are removing emoji and removing test from our library
smptw = '@user #white #supremacists want everyone to see the new 🐦 #birdsâ€ #movie â€ and hereâ€s why'
smptw = remove_emoji(smptw)
smptw = clean_text(smptw)
print(smptw)
```

```
user white supremacists want everyone birdsâ€ movie hereâ€s
```

```
data['tweet'] = data['tweet'].apply(remove_emoji)
data['tweet'] = data['tweet'].apply(clean_text)
data['num_words_text'] = data['tweet'].apply(lambda x : len(str(x).split()))
```

```
train_data, val_data = train_test_split(data, test_size=0.2)
train_data.reset_index(drop=True, inplace=True)
val_data.reset_index(drop=True, inplace=True)
```

```
# Now we are training the data set to count the words
test_data = val_data
print("==== Train Data ====")
print(train_data['label'].value_counts())
print(len(train_data))
```

```
print("==== Test Data ====")
print(test_data['label'].value_counts())
print(len(test_data))
```

```
==== Train Data ====
label
```

```

0    23747
1     1822
Name: count, dtype: int64
25569
==== Test Data ====
label
0     5973
1      420
Name: count, dtype: int64
6393

```

```

X_train, X_valid, y_train, y_valid = train_test_split(train_data['tweet'].tolist(), train_data['label'].tolist(), test_size=0.2, stratify=train_data['label'].tolist())
print("Train Data len: ", len(X_train))
print("Class distribution: ", Counter(y_train))
print("Validation Data len: ", len(X_valid))
print("Class distribution: ", Counter(y_valid))

```

```

↗ Train Data len: 20455
Class distribution: Counter({0: 18997, 1: 1458})
Validation Data len: 5114
Class distribution: Counter({0: 4750, 1: 364})

```

```
X_train[5]
```

```
↗ 'having them days today feelingdown'
```

```
num_words=50000
```

```

tokenizer = Tokenizer(num_words=num_words, oov_token="<UNK>")
tokenizer.fit_on_texts(X_train)

```

```

# Now we are making an array to do the filteration
from keras.preprocessing.sequence import pad_sequences

```

```

# Tokenize the sequences
tokenized_train = tokenizer.texts_to_sequences(X_train)
tokenized_valid = tokenizer.texts_to_sequences(X_valid)
tokenized_test = tokenizer.texts_to_sequences(test_data['tweet'].tolist())

```

```

# Pad the sequences to the same length
max_len = max(max(len(seq) for seq in tokenized_train),
               max(len(seq) for seq in tokenized_valid),
               max(len(seq) for seq in tokenized_test))

```

```

x_train = pad_sequences(tokenized_train, maxlen=max_len)
x_valid = pad_sequences(tokenized_valid, maxlen=max_len)
x_test = pad_sequences(tokenized_test, maxlen=max_len)

```

```

train_labels = np.asarray(y_train)
valid_labels = np.asarray(y_valid)
test_labels = np.asarray(test_data['label'].tolist())

```

```

print("Train data: ", len(x_train))
print("Validation data: ", len(x_valid))
print("Test data: ", len(x_test))

```

```

#Tensorflow dataset
train_ds = tf.data.Dataset.from_tensor_slices((x_train, train_labels))
valid_ds = tf.data.Dataset.from_tensor_slices((x_valid, valid_labels))
test_ds = tf.data.Dataset.from_tensor_slices((x_test, test_labels))

```

```

↗ Train data: 20455
Validation data: 5114
Test data: 6393

```

```

# Assuming sequences_train is a list of sequences
maxlen = max(len(seq) for seq in sequences_train)

```

```

# Now you can set the sequence_length
sequence_length = maxlen

```

```
max_features = 50000
embedding_dim = 16
sequence_length=maxlen

model = tf.keras.Sequential()
model.add(tf.keras.layers.Embedding(max_features + 1, embedding_dim, input_length=sequence_length, embeddings_regularizer=regularizers.l2(0.001)))
model.add(tf.keras.layers.Dropout(0.4))
model.add(tf.keras.layers.LSTM(embedding_dim, dropout=0.2, recurrent_dropout=0.2, return_sequences=True, kernel_regularizer=regularizers.l2(0.001)))
model.add(tf.keras.layers.Flatten())
model.add(tf.keras.layers.Dense(512, activation='relu', kernel_regularizer=regularizers.l2(0.001), bias_regularizer=regularizers.l2(0.001)))
model.add(tf.keras.layers.Dropout(0.4))
model.add(tf.keras.layers.Dense(8, activation='relu', kernel_regularizer=regularizers.l2(0.001), bias_regularizer=regularizers.l2(0.001)))
model.add(tf.keras.layers.Dropout(0.4))
model.add(tf.keras.layers.Dense(1, activation='sigmoid'))

model.summary()
```

↗ Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
embedding (Embedding)	(None, 20, 16)	800016
dropout (Dropout)	(None, 20, 16)	0
lstm (LSTM)	(None, 20, 16)	2112
flatten (Flatten)	(None, 320)	0
dense (Dense)	(None, 512)	164352
dropout_1 (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 8)	4104
dropout_2 (Dropout)	(None, 8)	0
dense_2 (Dense)	(None, 1)	9
=====		
Total params: 970593 (3.70 MB)		
Trainable params: 970593 (3.70 MB)		
Non-trainable params: 0 (0.00 Byte)		

```
model.compile(loss=tf.keras.losses.BinaryCrossentropy(), optimizer=tf.keras.optimizers.Adam(1e-3), metrics=[tf.keras.metrics.BinaryAccuracy()])

epochs=100
history=model.fit(train_ds.shuffle(5000).batch(1024), epochs=epochs, validation_data=valid_ds.batch(1024), verbose=1)
```



```

Epoch 77/100
20/20 [=====] - 2s 113ms/step - loss: 0.1685 - binary_accuracy: 0.9717 - val_loss: 0.2753 - val_binary_a
Epoch 78/100
20/20 [=====] - 3s 137ms/step - loss: 0.1701 - binary_accuracy: 0.9687 - val_loss: 0.3169 - val_binary_a
Epoch 79/100
20/20 [=====] - 3s 171ms/step - loss: 0.1719 - binary_accuracy: 0.9702 - val_loss: 0.2846 - val_binary_a
Epoch 80/100
20/20 [=====] - 2s 109ms/step - loss: 0.1743 - binary_accuracy: 0.9674 - val_loss: 0.2701 - val_binary_a
Epoch 81/100
20/20 [=====] - 2s 106ms/step - loss: 0.1653 - binary_accuracy: 0.9722 - val_loss: 0.2882 - val_binary_a
Epoch 82/100
20/20 [=====] - 2s 111ms/step - loss: 0.1635 - binary_accuracy: 0.9720 - val_loss: 0.2768 - val_binary_a
Epoch 83/100
20/20 [=====] - 2s 104ms/step - loss: 0.1694 - binary_accuracy: 0.9698 - val_loss: 0.2418 - val_binary_a
Epoch 84/100
20/20 [=====] - 4s 179ms/step - loss: 0.1737 - binary_accuracy: 0.9691 - val_loss: 0.2656 - val_binary_a
Epoch 85/100
20/20 [=====] - 2s 110ms/step - loss: 0.1657 - binary_accuracy: 0.9735 - val_loss: 0.2833 - val_binary_a
Epoch 86/100
20/20 [=====] - 2s 108ms/step - loss: 0.1617 - binary_accuracy: 0.9722 - val_loss: 0.2769 - val_binary_a
Epoch 87/100

```

```

predictions = model.predict(x_test)
print(predictions)

```

```

➡ 200/200 [=====] - 1s 5ms/step
[[1.1709050e-07]
 [3.1142645e-13]
 [1.0686236e-12]
 ...
 [2.5619832e-07]
 [2.7241500e-09]
 [1.2160532e-02]]

```

```


def display_training_curves(training, validation, title, subplot):
    _, ax = plt.subplots(figsize=(10,5), facecolor='#F0F0F0')
    plt.tight_layout()
    ax.set_facecolor('#F8F8F8')
    ax.plot(training)
    ax.plot(validation)
    ax.set_title('model ' + title)
    ax.set_ylabel(title)
    #ax.set_ylim(0.28,1.05)
    ax.set_xlabel('epochs')
    ax.legend(['train', 'valid.'])

```


```

display_training_curves(
    history.history['loss'],
    history.history['val_loss'],
    'loss', 211)
display_training_curves(
    history.history['binary_accuracy'],
    history.history['val_binary_accuracy'],
    'accuracy', 212)

```



```
_, (ax1, ax2) = plt.subplots(1, 2, figsize=(15,5))
ax1.scatter(predictions, range(0, len(predictions)), alpha=0.2)
ax1.set_title("Distributions")
ax2 = sns.distplot(predictions)
```



```
final_test_df = pd.read_csv('/content/test_tweets_anuFYb8.csv')
fctest = final_test_df.copy()
fctest.drop(columns=['id'], axis=1, inplace=True)

fctest['tweet'] = fctest['tweet'].apply(remove_emoji)
fctest['tweet'] = fctest['tweet'].apply(clean_text)

# Convert the list of lists to a NumPy array with object dtype to handle varying lengths
f_test = np.array(tokenizer.texts_to_sequences(fctest['tweet'].tolist()), dtype=object)
```

```
# Pad the sequences after converting to a NumPy array
f_test = pad_sequences(f_test, padding='post', maxlen=maxlen)
```

```
display(f_test)
```

```
array([[14136,      1, 9280, ...,      0,      0,      0],
       [      2, 186,      1, ...,      0,      0,      0],
       [ 496, 666, 519, ...,      0,      0,      0],
       ...,
       [ 749,      1, 17, ...,      0,      0,      0],
       [      9, 32, 290, ...,      0,      0,      0],
       [ 307, 644, 121, ...,      0,      0,      0]], dtype=int32)
```

```
predictions_f_test = model.predict(f_test)
```

```
538/538 [=====] - 2s 4ms/step
```

```
_, (ax1, ax2) = plt.subplots(1, 2, figsize=(15,5))
ax1.scatter(predictions_f_test, fctest.index, alpha=0.2)
ax1.set_title("Distributions")
ax2 = sns.distplot(predictions_f_test)
```